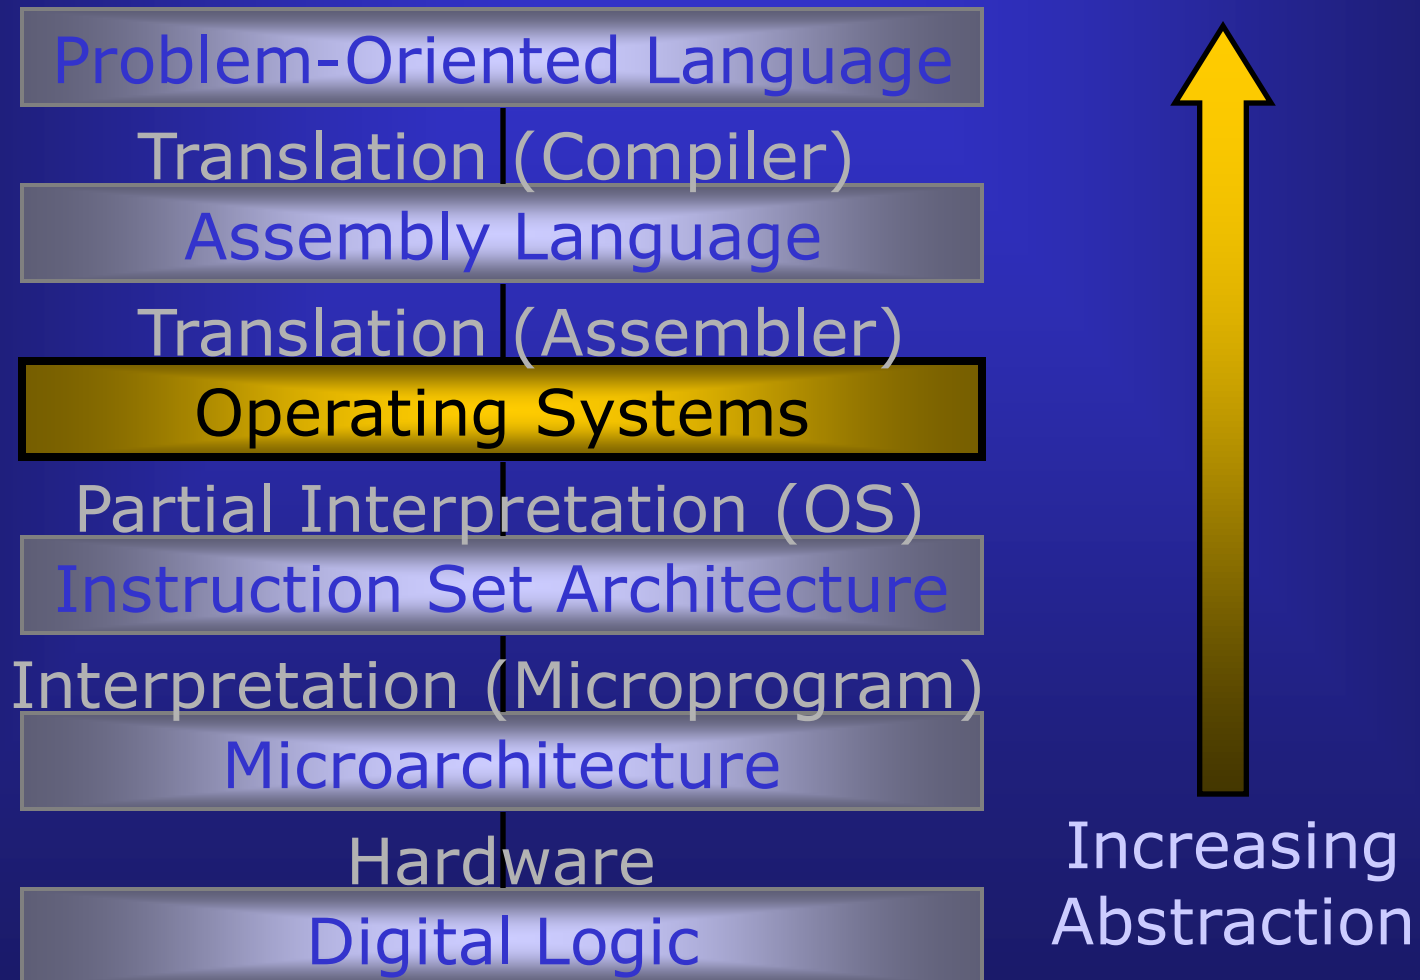
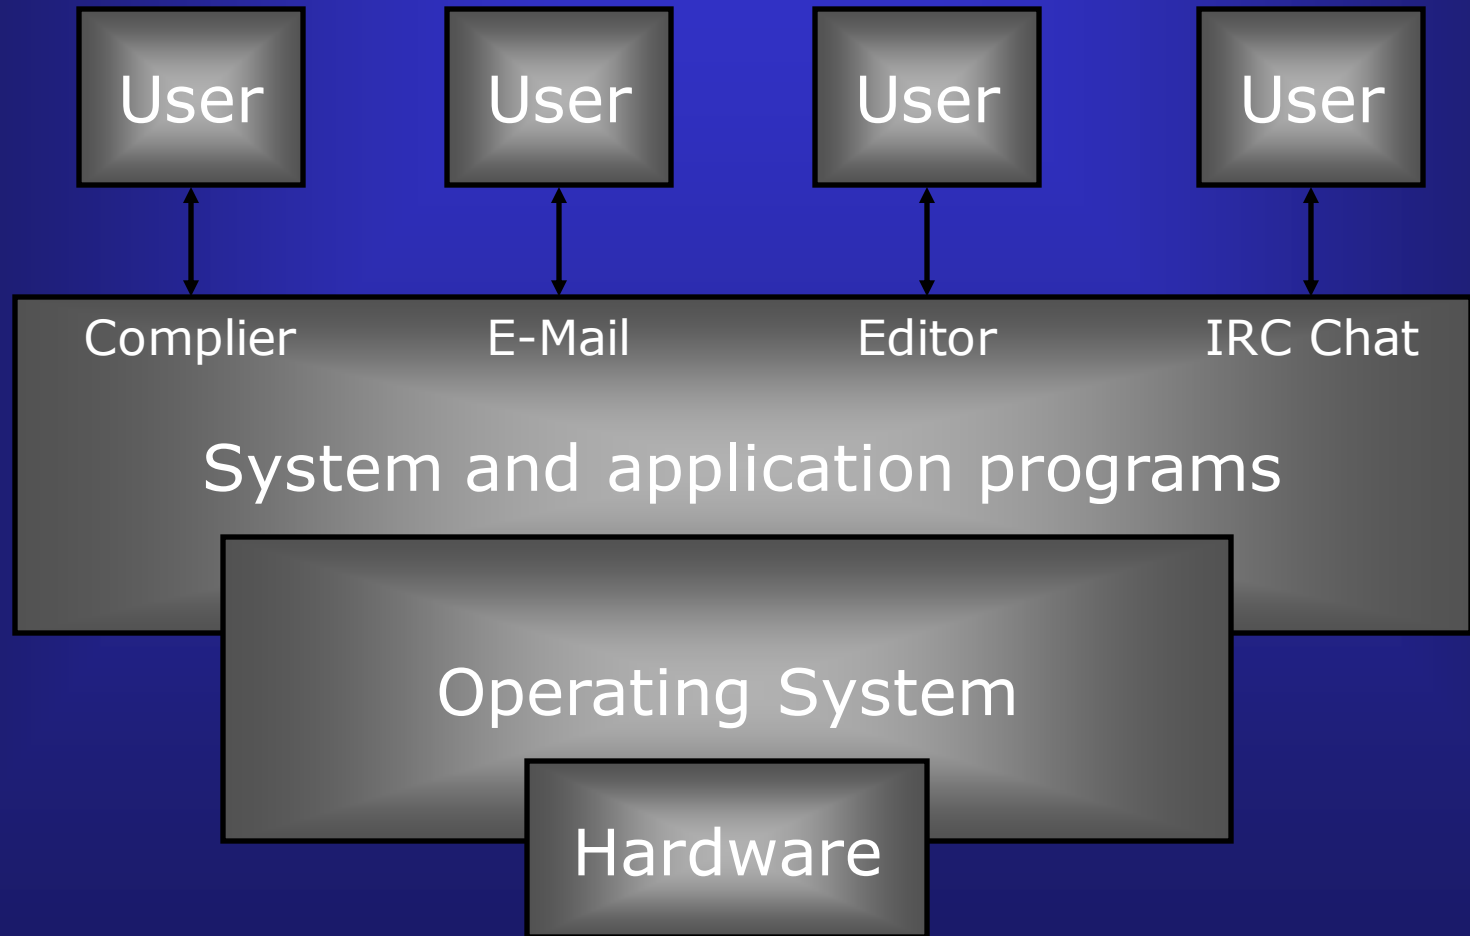


Multilevel Machines



The Operating System Layer



Operating System Tasks

- ■ Process Management
- ■ Memory Management
- ■ File-System Management
- ■ Communications
- ■ Error Detection
- ■ Accounting
- ■ Protection

Protection: The Need for Hardware Support

- ■ Need to ensure that no program can interfere with any other
 - ■ Could overwrite other programs
 - ■ Need to protect the OS in particular
- ■ Graceful recovery from errors detected by the hardware
 - ■ E.g. illegal instructions, divide-by-zero...
- ■ Hardware invokes the OS
 - ■ *Trap*: a “software interrupt”
 - ■ CPU transfers control to an address stored in a trap register

Dual-Mode Operation

- ■ At least two separate modes, implemented using a *mode bit*
 - ■ User mode (mode bit = 0)
 - ■ Supervisor mode (mode bit = 1)
- ■ OS runs in supervisor mode
- ■ User mode is intended to restrict access to the machine
- ■ Also need *privileged instructions*
 - ■ Executed only in supervisor mode
 - ■ Cause a trap if executed in user mode
 - ■ E.g access to the mode bit

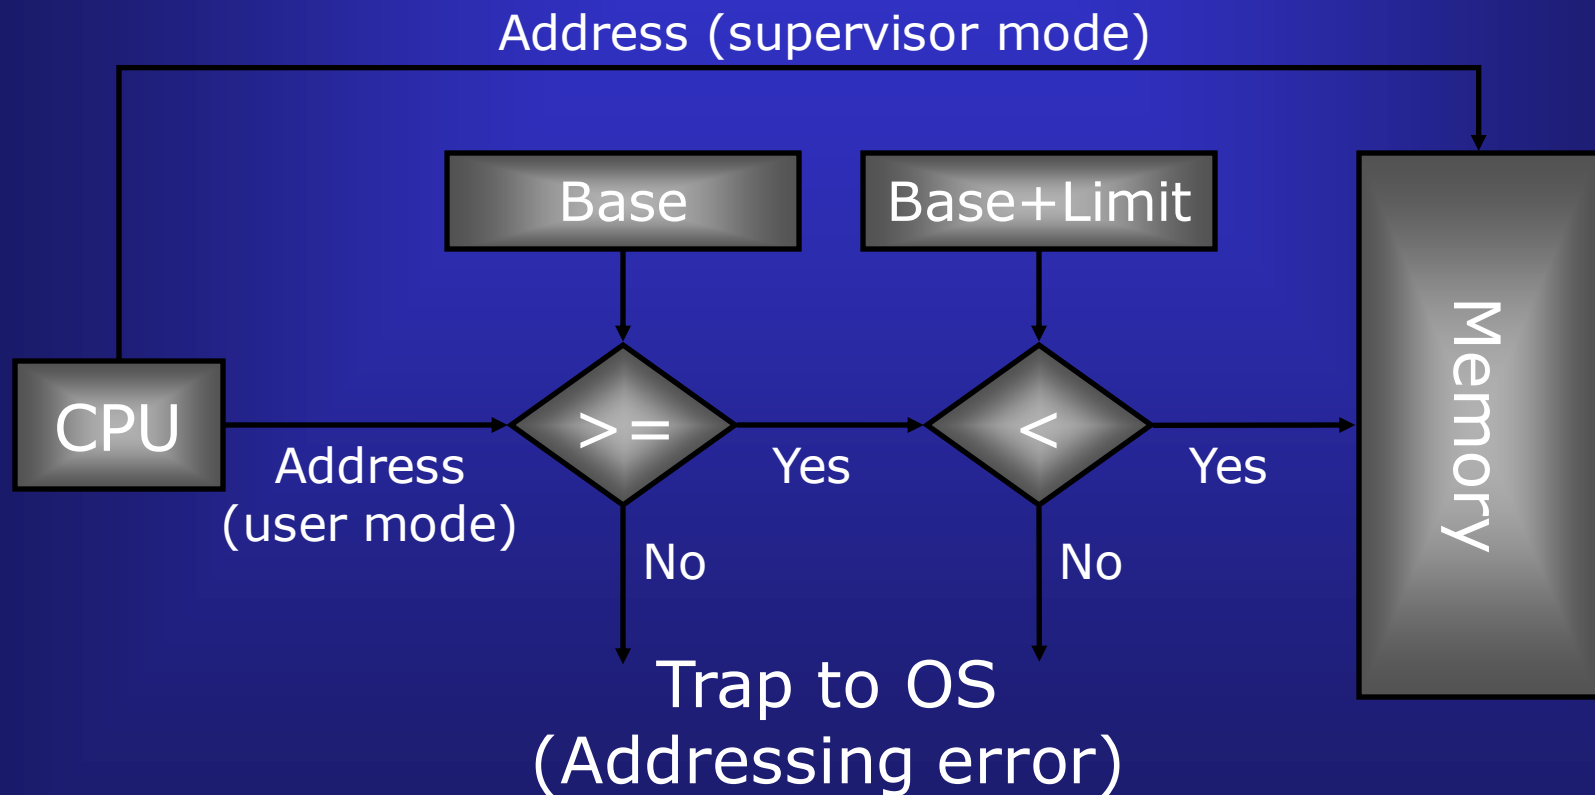
Privileged Instructions: I/O Protection

- ■ User program could issue illegal I/O instructions
 - ■ Write to a non-existence device
 - ■ Read more data than a disk holds
- ■ Solution: make *all* I/O instructions privileged
- ■ User programs can then only perform I/O by requesting it through the OS
- ■ OS retains control over user programs

Memory Protection

- ■ Need to limit access to main memory
 - ■ Prevent modification of OS and other programs
- ■ Give each program a range of legal addresses
 - ■ But need to compare *every* address issued by the user's program
- ■ *Base* and *Limit* registers
 - ■ Implemented in hardware
 - ■ Privileged instructions used to set them
- ■ Supervisor mode has unlimited access

Base and Limit Registers



CPU Protection

- Need to prevent programs from hogging the CPU
 - Infinite loops
 - Waiting for non-existent resources
- Use a *timer* to control program execution
- When timer expires, switch to OS
 - *Context switch*
- This is the idea of *time-slicing*
 - Each program runs for a few msec

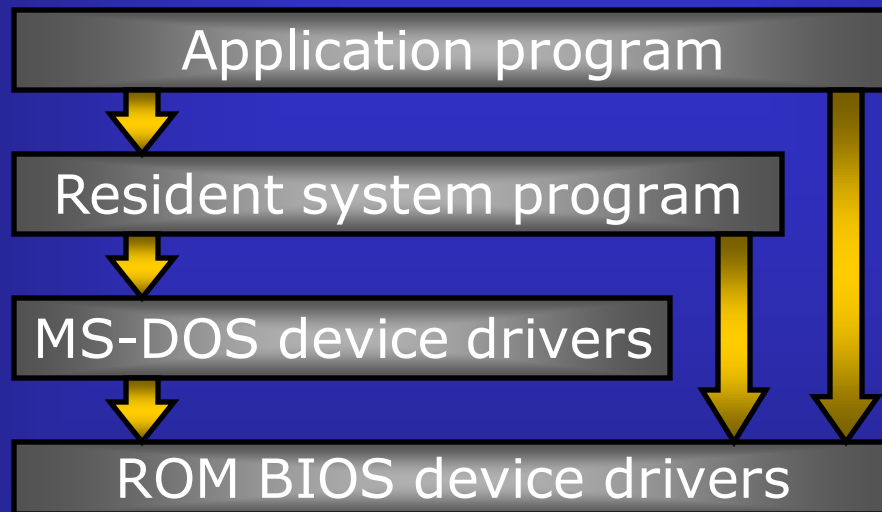
System Calls (the *API*)

- Provide the interface between a process and the OS
 - Process manipulation
 - File and I/O manipulation
 - Information maintenance
 - Communications
- Ask the OS to perform some task on the processes behalf
- Some languages (C, C++, perl,...) allow system calls to be made directly
 - Java hides all system calls

System Programs

- Provide an environment for program development and execution
- Interfaces to system calls
 - Status information (date, time, space...)
 - File modification/manipulation (file creation, deletion, renaming, various editors,...)
 - Programming support (compilers, linkers, debuggers, assemblers, interpreters...)
 - Program loading and execution (command interpreters, graphical interfaces...)
 - Communications (email, ftp, telnet...)

Simple Structure: MS-DOS

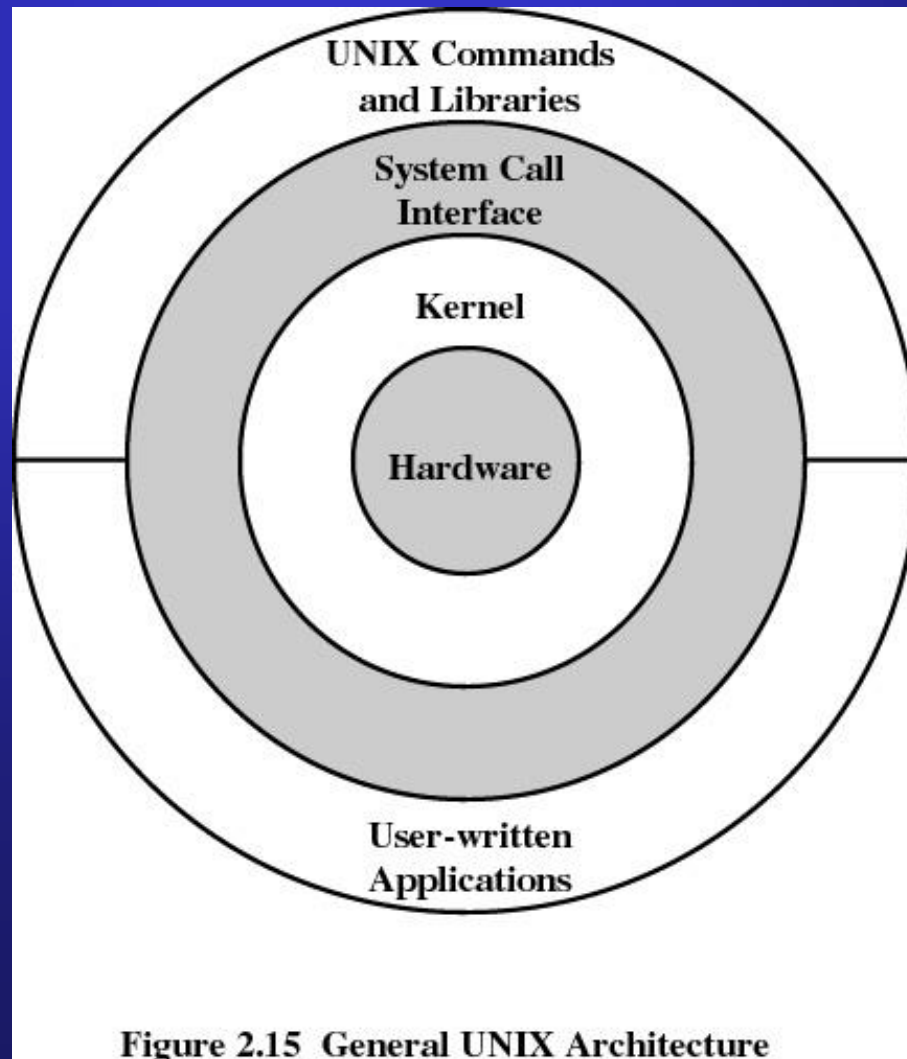


- ■ No clear separation of functionality
- ■ No dual-mode, so applications can access hardware directly
 - ■ Fast, but potentially unsafe

PORTABLE OS

1. Problems:
 - a) Architecture and OS have strong ties due to which computer manufactures face difficulties in marketing new Architectures as OS development take substantial time
 - b) Develop new set of skills for each OS and M/c
2. Hence portability is required and standardization is required from both user and architecture point of view
3. Two components
 - a) Policy : governs the use of resource, RR policy
 - b) Mechanisms : implements the policies, RR implementation
4. Portable OS is structured so as to segregate the policy from the mechanism. This makes the policy implementation of a m/c on which the OS is to be used
5. Mechanism : Kernel of OS to be rewritten for the new m/c
6. UNIX supports portability, layered OS
7. Now work is going on for micro-kernel design to introduce more portability.13

UNIX

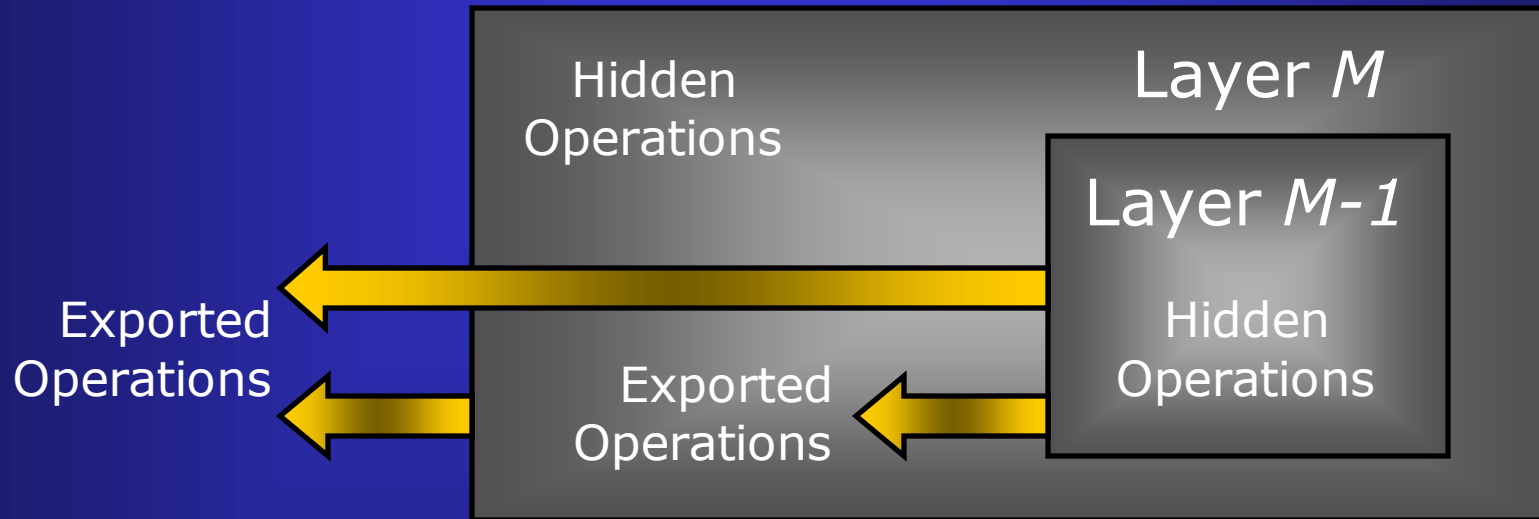


Kernel + System Programs: UNIX

Users		
Shells and commands Compilers and interpreters System libraries		
System-call interface to kernel		
Terminal handling Character I/O Terminal drivers	File system Block I/O Disk/tape drivers	CPU scheduling Page replacement Demand paging Virtual Memory
Kernel interface to hardware		
Terminal controllers Terminals	Device controllers Disks and tapes	Memory controllers Physical memory


- ■ Easy to port to different hardware
- ■ But the kernel is very complex

Layered Structure: OS/2



- ■ Bottom Layer = hardware
- ■ Top layer = user interface
- ■ Modularity: very easy to debug
 - ■ Difficult to assign appropriate functions to each layer
 - ■ Layers add overhead, so can be inefficient

Microkernels: NT



- ■ All non-essential parts of kernel are user-level programs
 - ■ Reliability and security
- ■ Message-passing communication
 - ■ Client/Server architecture
- ■ Easy to extend OS
 - ■ Don't need to modify kernel

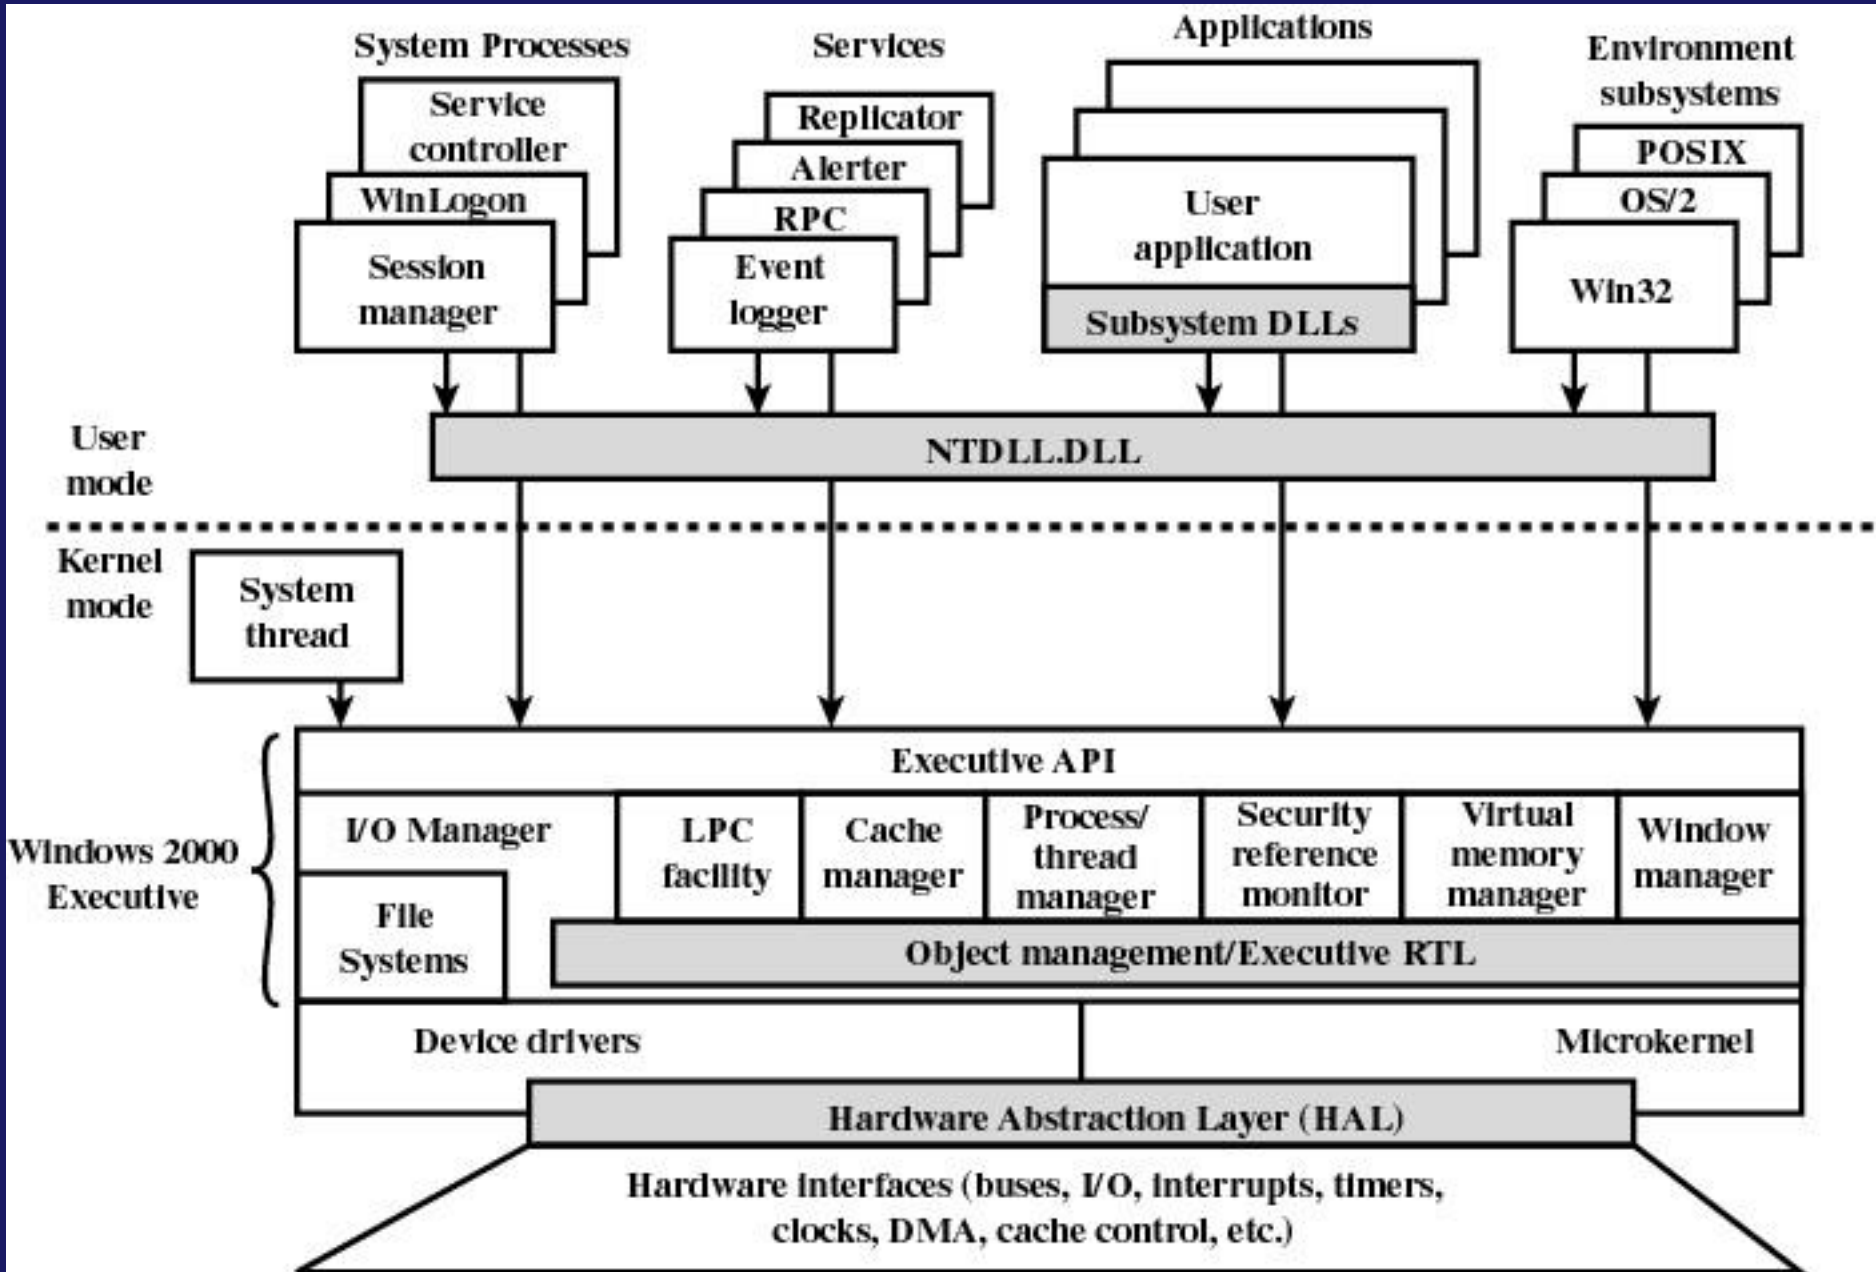
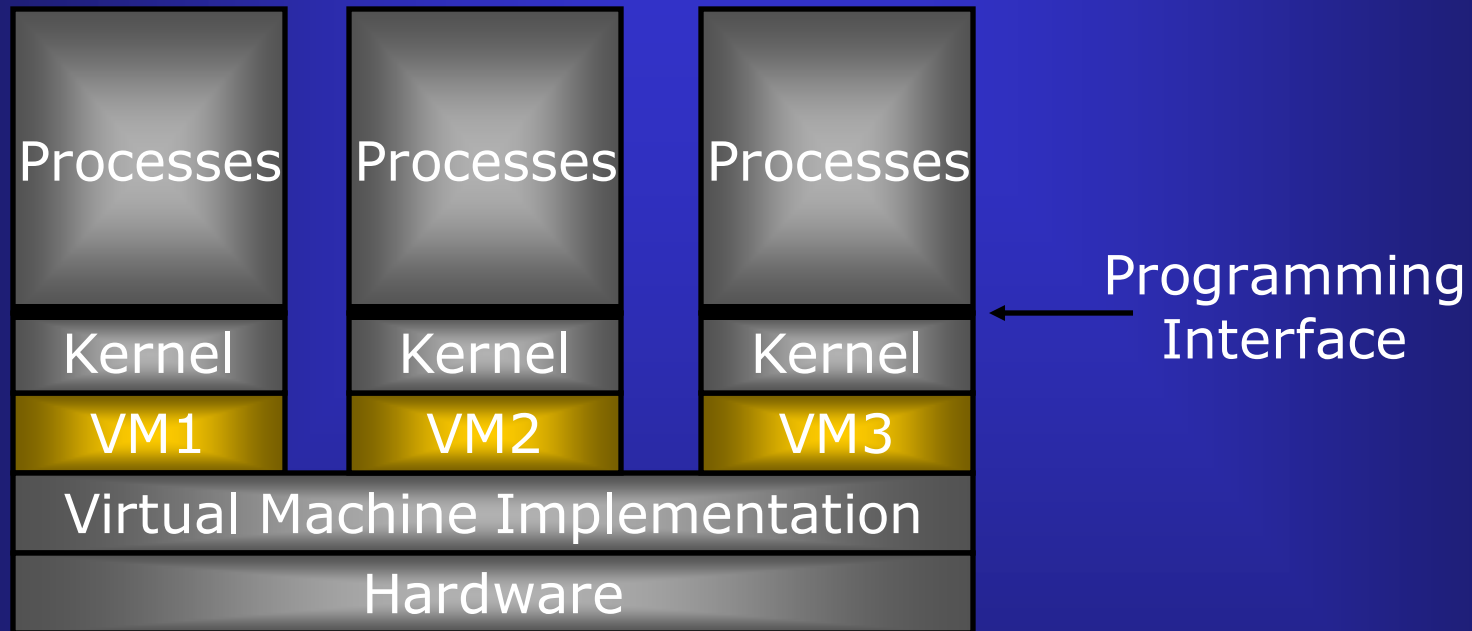


Figure 2.13 Windows 2000 Architecture

Virtual Machines: IBM




- Provide processes with a virtual copy of underlying hardware
 - Instruction interpretation
- Users have their own virtual machine
 - Difficult to provide *exact* duplicate of hardware

Processes and Threads

We will see:


- ■ What a process is, and what it does
- ■ The idea of concurrent execution
- ■ What a thread is
- ■ How processes are implemented in the operating system
- ■ How the operating system decides which job to run

The Process Concept



- ■ A *program* is a passive entity
 - ■ Just a sequence of instructions
- ■ A *process* is an active entity
 - ■ The program in execution
- ■ A *processor* is an agent which executes processes

Concurrency



■ ■ *Concurrent processes*

- ■ Activation of more than one process

- ■ Apparent concurrency can be achieved by switching a processor between several processes

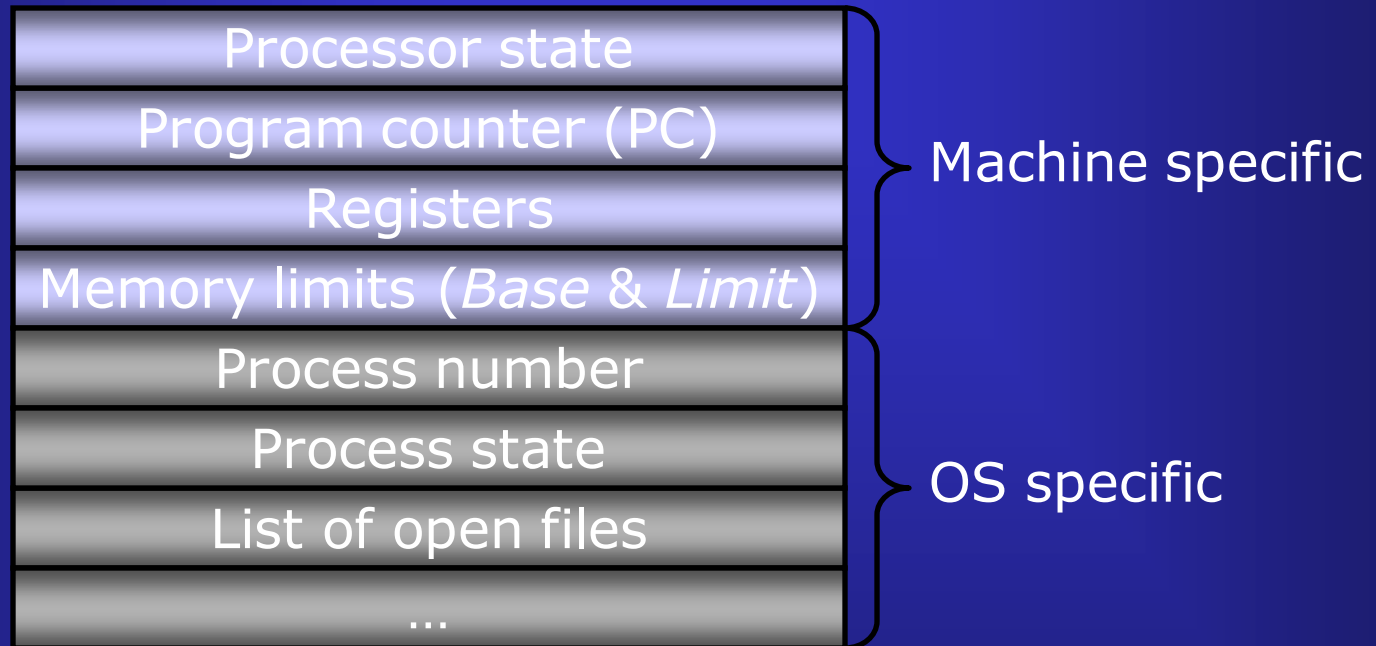
- ■ *Instruction interleaving*

Processes can Share Code



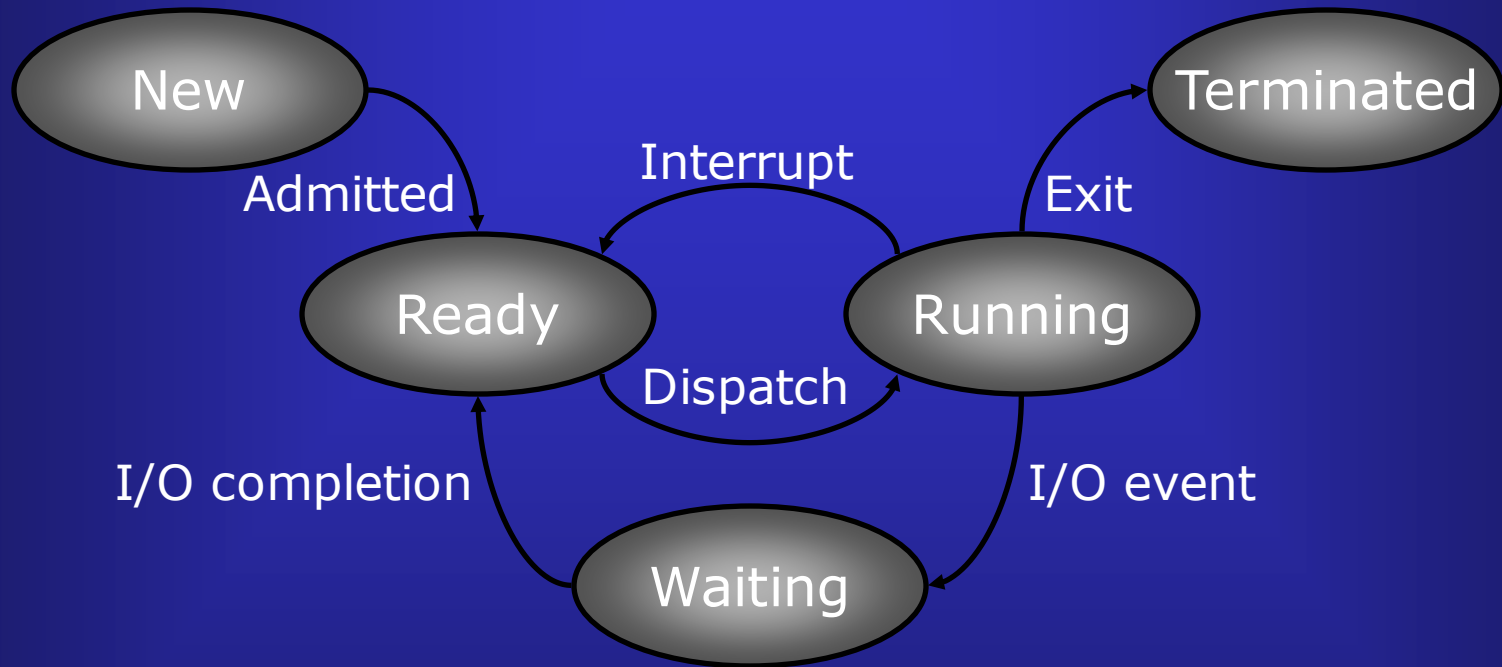
- Two concurrent instances of a text-editor
- Each uses the same program, but has different data and state information

Process Control Block



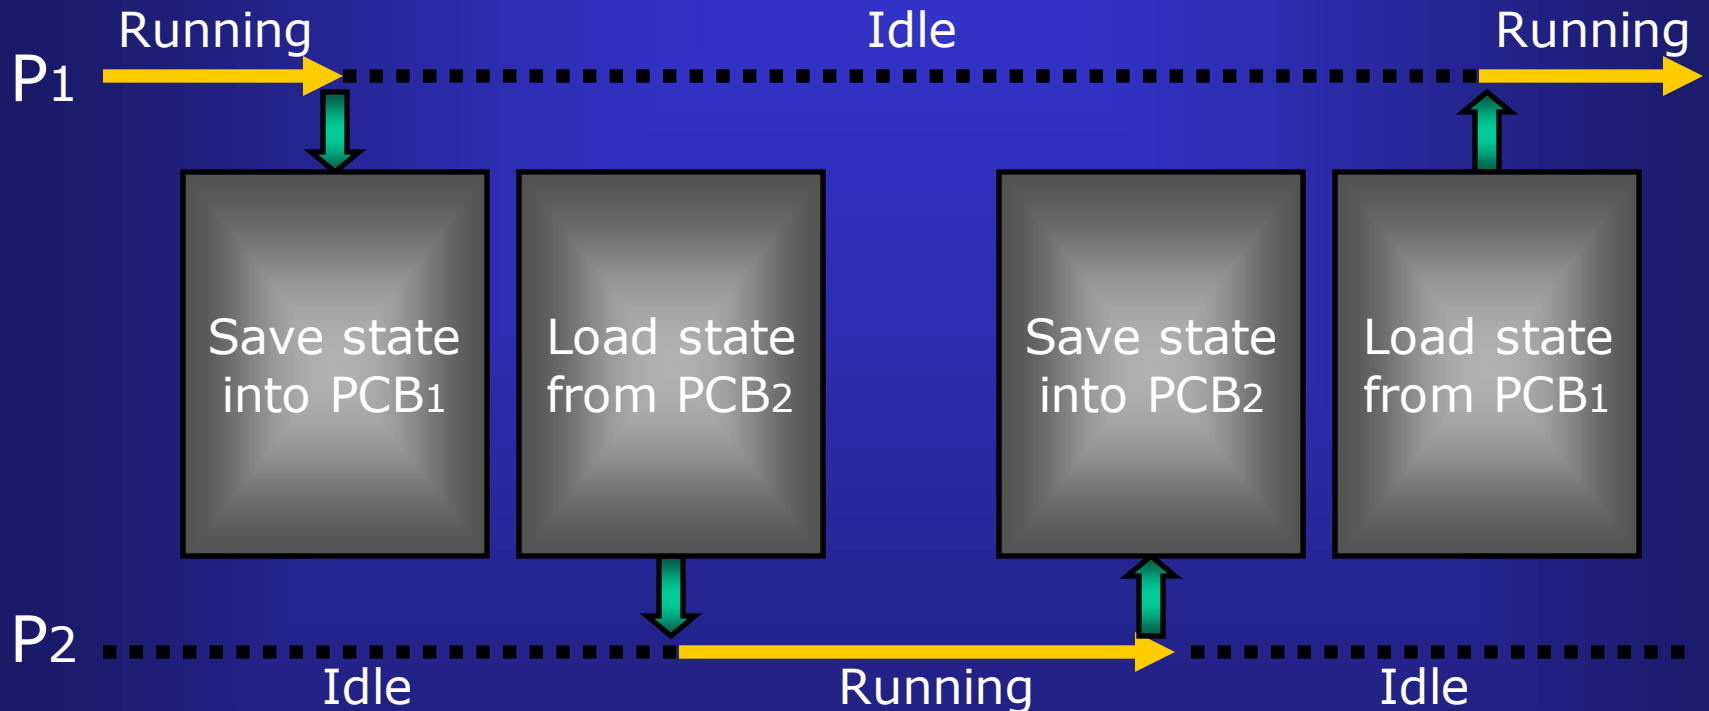
- ■ Representation of a process in the OS
- ■ Contains all the data about the process
 - ■ One PCB per process

Process State



- ■ As a process executes, it changes *state*
 - ■ Only one *running* process per CPU
 - ■ Many *ready* or *waiting*

Context Switching




- Save state of current process
- Load state of next process
- The CPU does no useful work while switching

Threads

- Context switches take time
 - As OS becomes more complex, the amount of data associated with a process increases
- A *thread* is a lightweight process
 - One process can have many threads
 - Separate machine state
 - Share OS resources (code, files,...)
- Kernel threads (NT, Digital UNIX)
 - Run as separate entities in OS
- User threads (thread libraries)
 - Run under a single process

Swapping



- Swapped-out processes are kept on disk
- Aim to swap out *non-interactive* processes
- Improves the performance of interactive processes

Process Scheduling


- Aim to have a process running at all times
 - Time-sharing
 - Maximise use of the CPU
- The *ready queue*
 - List of processes which are ready to run
- Processes remain in the ready queue until they are *dispatched*
- The *scheduler* decides which process is to run next
 - How it does this affects the whole system

Priority Pre-emption



- Each process is assigned a priority
- Processes run until they suspend themselves or are interrupted
- If an interrupt makes a higher-priority process ready, then run it instead

Run-To-Completion



- ■ Process runs until it suspends itself
- ■ After an interrupt, the same process continues to run




Interrupt Latency

- Time between an interrupt occurring and a process being run in response to that interrupt
- $\text{Interrupt Latency} = \text{Interrupt lock-out time} + \text{Time to service interrupt} + \text{Time to schedule process}$
- Advantages of priority pre-emption
 - Interrupt latency is constant
 - Can have low-priority processes which do not interfere with response times

Interrupt Latency with Natural Break

Ready queue =  

Interrupt puts P3 onto ready queue

Ready queue =   

P1 runs for 2 msec, P2 runs for 3 msec

Interrupt latency with natural break = 5 msec


■ ■ Advantages

- ■ Efficient implementation of mutual exclusion

■ ■ Disadvantages

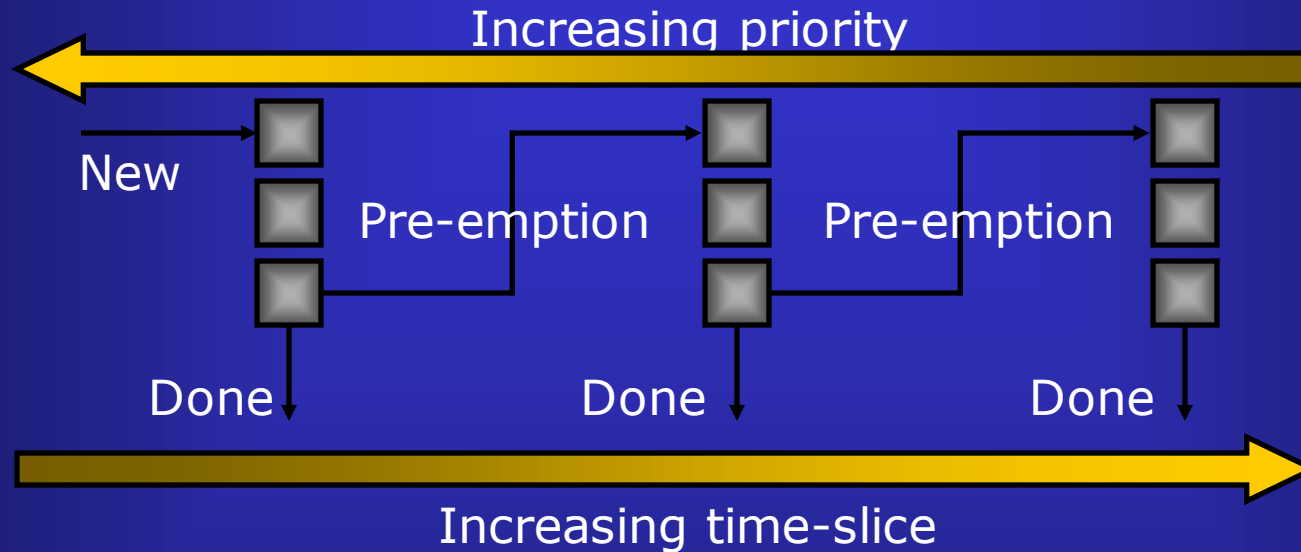
- ■ User processes must be well-behaved
- ■ Response time depends on user programs

Time-Sliced Scheduling



- Process runs until its *time-slice* expires or until it suspends itself
- Ensures that CPU is shared fairly between users
- Interrupt latency is the size of the time-slice

Multi-Level Queues



- ■ Dynamic priority assignment
- ■ Favours interactive jobs
 - ■ Get a good response time
- ■ Used in UNIX

Summary

- ■ The fundamental unit of computation is a *process*
- ■ Managing processes one of the most important jobs of an operating system
- ■ The operating system uses a scheduling strategy to decide which process to run when

Types of Process Interaction

■ ■ Sharing

- ■ Two processes sharing an object
- ■ Require *mutually exclusive* access to prevent interference

■ ■ Synchronisation

- ■ One process triggering another

■ ■ Communication

- ■ One process passing a message to another

■ ■ These are all closely related